Estimating error rates for bullet comparisons in forensic science

Yawei Ge[,1], Heike Hofmann[1]

[a]*Department, Street, City, State, Zip*
[b]*Department, Street, City, State, Zip*

Abstract

This is the abstract.

It consists of two paragraphs.

Text based on elsarticle sample manuscript, see http://www.elsevier.com/ author-schemas/latex-instructions#elsarticle

Introduction and literature review Part1

The firearm examiners are focusing on the same source and different source problems of bullets and cartridge cases which serve as important forensic evidence in the court/to the jurors. XX this is quite vague - could you try to find a number in the literature on how much evidence in trials is related to firearms or ballistic evidence? The subjectivity and lack of quantified error rates in the traditional forensic process is called to be reduced or complemented by more objective methods (**???**; National Research Council, 2009). Some XX 'some' is too vague - instead cite literature - i.e. something like: in response xxx, yyy, zzz, ... have introduced automatic matching algorithms ... automatic matching algorithms are developed which usually return a similarity score to quantify the similarity or the probability to be an actual match for a certain comparison. However, this raises questions about how to interpret the reported scoresand how these scores are distributed. XXX more details before this thus - why is it XXX it? write out your reference not clear to draw inference? Thus, it XXX it? write out your reference is not all clear how to conduct inference based on these similarity scores. Song et al. (2018) proposed binomial and beta-binomial for the number of matched cells of the CMC method for cartridge case comparisons. Therefore, Song et al. (2018) provides a way to quantify the theoretical error rate of the algorithm. However, for the bullets LEA comparisons, quantitative measurements for the theoretical error rate have not been established. In this paper, we will evaluate the possible models/distributions for the LEA comparisons scores generated by the random forest proposed by Hare et al. (2016). And then, we will also

---

evaluate the error rates based on the estimated distribution for the automatic matching algorithm.

In section 2, we will discuss distributional forms of random forest scores produced by the automatic LEA matching algorithm proposed by Hare et al. (2016). In section 3, we will introduce the LAPD data set and the estimated distributions. In section 4, the theoretical error rates based on the distributions are discussed. In section 5, we evaluate the performance of the estimation, stability of the distribution within a changing sample size context. In section 6, we will conclude the discussion.

To combine the following introduction Part2

The subjectivity in the conventional forensic evidence identification processes according to the AFTE is called to be reduced or be complemented by more objective procedures (National Research Council, 2009). Don't use passive voice in the start The three possible conclusions XXX which conclusions are you refering to? you need to introduce that before you can refer to them for a comparison between a known pattern (from tests of the tools or firearms of the suspect) and an unknown pattern (extracted from the crime scene) according to AFTE (AFTE Criteria For Identification Committee, 1992) are identification, elimination and inconclusive. Besides, for some comparisons which are thought to be impossible to do any meaningful judgement, a forth decision will be made to exclude those cases. XXX Careful: exclusions are based on single pieces of evidence. Only if evidence is considered to carry information, a comparison is done. The conventional standards established by AFTE were usually carried out by experienced examiners. However, this process in nature is human based and subjective. the previous sentence explains the foundation, put it earlier. In recent years, this procedure was called to be changed. XXX repeat from first sentence. Especially with the improvement of the technology in measurement tools and higher computational power. (some more general discussions could be made, such as the power/usefulness of this kind of evidence, the current stage of those procedures, how to cooperate with human judgement)

The identification results of bullets sources are part of the important evidence to jurors. XXX are there numbers of how often firearms evidence is presented in court? When the gun fires a bullet, the barrel and the rifling will leave an impression on the bullet surface. For different barrels and rifling, the marks left on the bullets are believed to be unique (AFTE Criteria For Identification Committee (1992)). On the highest level, we can easily distinguish some class characteristics such as numbers of land engraved areas, types of bullets etc. When the bullet from the crime scene is different from the test bullets which are test fired from a suspect gun in terms of the class characteristic, the decision can be made easily as elimination.(what about sub-class characteristic) However, we can never make identification solely based on class characteristics. The bullet forensic researchers compare the individual characteristics, which are usually due to the imperfectness of guns or wear of guns, thus leaves unique marks on the bullets. The bullet examiner in the US labs will put two bullets under the

confocal microscope, and try to match the striae on each bullet visually. With technology improvement, researchers are able to get 3D scans of the bullets and develop more advanced comparison algorithms. Hare et al developed a random forest based automatic matching algorithm with some well used features on the Hamby data (Hare et al. (2016)). They developed a whole process to start with the 3D scan and go through a sequence of data manipulations to reach the final comparison. Vanderplas et al also applied this method to another three external data sets Hamby set 44, Phoenix PD, Houston FSC (Vanderplas et al. (2020)). According to their study, the random forest algorithm is the best while ccf is the most prominent feature. Hare et al extended this approach to include degraded cases and operator effect into consideration (Hare et al. (2017)). Rice improved the groove engraved area identification algorithm by using the Robust LOESS (Rice et al. (2020)). Krishnan proposed Chumbley score method to match striae on the LEAs (Krishnan and Hofmann (2019)).

The automatic identification processes always produce a similarity score (usually in range 0 to 1) or a probability to be an actual match as results(references). This raises questions about how to interpret the reported probabilities and how these probabilities are distributed. In practice, the default cutoff of similarity score above which to be classified as matches is usually 0.5 with intuitive interpretation. However, without knowing how the similarity score is distributed and the classification errors associated with it, a sensible decision is hard to make. Song et al used a beta-binomial distribution to model the CMC cartridge cases classification results (Song et al. (2018)). Garton et al illustrated the use of score based likelihood in evaluating the strength of forensic evidence (Garton et al. (2020)).

In this paper, we focused on the bullet LEAs, and the results returned from the automatic matching algorithm of Hare et al (Hare et al. (2016)). We make use of the LAPD bullets data base (reference or explanation, in a separate section), which provided a large number of samples making the modeling process and evaluating the model possible to us. (The priciple outcome)

This needs a more general introduction to firearms examinations to motivate the problem

things the intro needs to cover:

- define the same-source problem

- cross-correlation function

  Based on the cross-correlation function we produced through...,

- discuss approaches in the literature: (NIST paper) on modelling Beta distribution

- how does this paper expand on the knowledge?

The distributional forms of the similarity scores

The quantitative methods used to objectively measure the similarity between LEAs reports various quantities, such as counts, correlations, distances, prob-

abilities and more generally similarity scores [references]. To understand how those quantities reflect the strength of evidence and to study the underlying the error rates in making decisions based on those quantities, distributional forms are usually set up [references]. Particularly, we are focusing on the similarity scores which range from 0 to 1 and the probabilities reported as the likelihood of an actual match. The similarity scores reported in the forensic researches are usually categorized into two categories. One is that the compared LEAs or other forensic evidences are actual matches, the other is that the compared LEAs or other forensic evidences are not matches. We name the former as known matches (KM) and the latter as known non-matches (KNM), and the corresponding distributions are named as KM distributions and KNM distributions as in Figure . . . for example. However, this kind of classifications can be investigated only in the lab environment where the ground truth is known. When we make any decisions based on any quantitative measurement, we are actually making a distinguish between those two potential distributions. The strength of any identification process is also measured by the disparities of those distributions. However, in practice, we can hardly discriminate those two distributions entirely, thus, we are never 100% sure which distribution the observed score comes from. This is where the identification error raises.

For the purpose of illustration, we have a look at the example (could use Hamby and other sets, not necessarily LAPD) in Figure . . . . Those are RF scores generated from the automatic matching algorithm [references]. The similarity scores from the RF can be explained as probabilities that calculated through the algorithm based on the trained model to quantify the likelihood that a pair of LEAs are actually a match. Or we can think of the RF scores as general similarity scores that quantify the similarities. As the name indicated, the higher RF scores imply higher chances to be a match. For different combination of ammunition and firearms, the scores are distributed differently. It is expected that systematic differences exist there for different cases [references]. So, it is necessary to study the scores under controlled conditions, thus, thresholds for the scores to do classifications vary based on the changes of the underlying distributions.

We can see from the Figure . . . , which is a typical one we usually have for the scores, that the distributions of KM and KNM are apart for the majorities. In the bullet LEA comparison problems, we usually have a well separated bullet scores but for the land scores, there are usually some overlaps [Susan's paper]. We propose beta distributions for those scores. Because the beta distribution is a well-used one in statistics to describe a quantity from 0 to 1 which is usually a probability or proportion quantifying our knowledge for another distribution in Bayesian analysis. And it is very flexible to capture unimodal asymmetric shapes in 0 to 1. However, it may not be adequate to explain a heavy tail or even a second mode. Thus, we consider the beta mixture distribution which is a more complex distribution than the beta distribution as a special case. In the beta mixture distribution, we introduce a hierarchical structure with a prior probability to combine two beta distributions as one. The distributions are defined below.

The reported similarity scores are denoted as Y_{ij} for j^{th} LEA comparison within class i, where $i = 1$ is KM and $i = 2$ is KNM. Y_{ij}'s are considered independent and identically distributed within each class, i.e.

$$Y_{1j} \overset{iid}{\sim} Betamix(p_1, \mu_{11}, \phi_{11}, \mu_{12}, \phi_{12})$$

$$Y_{2j} \overset{iid}{\sim} Betamix(p_2, \mu_{21}, \phi_{21}, \mu_{22}, \phi_{22})$$

where μ_{ik} and ϕ_{ik} are distribution parameters for i^{th} class and k^{th} component, $k = 1$ or 2. And p_i is the prior probability for the first component in i^{th} class, thus, $1 - p_i$ is the prior probability for the second component in i^{th} class. Note that we are using the mean and precision parameterization of beta distributions which simplifies the math in calculation and is more intuitive (μ is the mean, and the variance is roughly proportional to the reciprocal of ϕ). It's equivalent to the usual α and β parameterization through the following transformation:

$$\mu = \frac{\alpha}{\alpha + \beta}$$

$$\phi = \alpha + \beta$$

LAPD data set and the estimated distributions

For the following sections of the paper, we will base our analysis on the LAPD data sets. This is a large data set of . . . It is the first time such a large data set available to the researchers, which makes it possible for a statistical analysis for the distribution of similarity scores. (More! And number of land comparisons for different cases)

We consider cross correlation functions (CCF) for land comparisons, which is produced by calculating the maximized CCF between two signatures extracted from a pair of bullet LEAs [refer to a paper, could be the random forest one or others].

The estimation was done using Nelder-Mead algorithm in R [reference]. This is a general purpose numeric method which works reasonably well for non-differentiable function. In our case, the non-differentiable objective function is the log likelihood function of the beta mixture distribution. Therefore, we finally found the maximum likelihood estimators (MLE). Instead of a simpler beta model, we start with the more complex two-component beta mixture model, see how well it fits the data and test if some components are necessary. We estimated beta mixture distributions for both KM and KNM as in Table 1 and Table 2 respectively. For the beta distributions, the KM distribution has a mean at 0.635 and the KNM distribution has a mean at 0.404. We can also see that for the two component beta mixture distributions for both KM and KNM, there are components with mean around 0.5 and the other components more separated from each other. This indicates the model successfully accounted for different situations of comparisons involving tank rash and random identification etc.

This is an ideal property we would like to see and make use in explaining the similarity scores. Another point about the estimates that worth mentioning is that the ϕ of the KM beta distribution is much smaller than any of the two components of the two-component KM beta mixture distribution. This indicates that the two component distribution could be better because of smaller variances for components.

Full Data KM Distribution Estimation						
Model	Component Prior Probability	μ	ϕ	logLik	p-value	BIC
3-comp	0.247	0.406	24.097			
	0.521	0.664	13.342	6278		-12479
	0.232	0.819	33.823			
2-comp	0.419	0.466	15.812	6243	0	-12438
	0.581	0.759	18.476			
Beta		0.635	6.529	5712	0	-11404

Table 1: Parameter estimations for the Beta distribution (1-component beta mixture distribution), 2-component and 3-component distributions for the KM CCF. "2-comp" refers to the 2-component beta mixture distribution, and the same for "3-comp". Column "logLik" is the maximized log likelihood for each distribution. Column "p-value" is the p-value for asymptotic likelihood ratio tests between the current model and one-step more complex model, where 0 indicates that we would reject the hypothesis that the current one is sufficient to decribe the data

Full Data KNM Distribution Estimation						
Model	Component Prior Probability	μ	ϕ	logLik	p-value	BIC
3-comp	0.650	0.345	35.017			
	0.314	0.490	24.979	78961		-157830
	0.032	0.667	18.454			
2-comp	0.674	0.358	39.908	78808	0	-157558
	0.336	0.494	13.324			
Beta		0.404	15.801	75349	0	-150675

Table 2: Parameter estimations for the Beta distribution (1-component beta mixture distribution), 2-component and 3-component distributions for the KNM CCF. "2-comp" refers to the 2-component beta mixture distribution, and the same for "3-comp". Column "logLik" is the maximized log likelihood for each distribution. Column "p-value" is the p-value for asymptotic likelihood ratio tests between the current model and one-step more complex model, where 0 indicates that we would reject the hypothesis that the current one is sufficient to decribe the data

The estimated beta distributions are shown in the Figure 1 and the estimated two-component beta mixture distribution are also shown in the Figure 2. As expected, the estimated distributions show the properties we desire for the

similarity scores. The majorities of the estimated distributions are apart from each other, while the minority part between the two distributions has some overlap. It's worth to note that both curves have heavy tails to the farther boundaries and the KM curve has the heavier tail compared to KNM. We can see the two-component distributions fit the data very well, while the single beta distributions are not as good as the two-component distributions for both KM and KNM. The single beta distribution for KM clearly failed to capture the potential second mode of the histogram indicating it is not sufficient. The single beta distribution for KNM is not too bad but still failed to capture the distributional information for some parts. This indicates that even though beta distributions are flexible for variables in 0 to 1, they are still restricted too much to be able to describe the cases here. The two-component beta mixture distributions look more promising.


Figure 1: Estimated beta distributions for full data

Figure 2: Estimated two-component beta mixture distributions for full data

Three candidate distributions for each of KM and KNM are considered. For the increasing of complexity, we have single beta, two-component beta mixture

and three-component beta mixture distributions. And they are also nested in that order. Naturally, we first look at the maximized log-likelihood of each model, and we can do asymptotic log-likelihood ratio chi-square tests for single beta against two-component beta mixture distributions, and two-component against three-component beta mixture distributions. The p-values of those test are shown in the Table 1 and Table 2 as the column "p-value". Surprisingly (or not), we found all those p-values are 0 which strongly suggests a more complex model when there is one. However, considering the size of the data we are using to fit these models, we can expect the statistical significance will be easily achieved since any small difference of the sufficiency will be detected. So we have to take the sample size effect into consideration. As Bayesian information criterion (BIC) is a well used criterion which takes the model complexity, sufficiency and the sample size (by a log function) into account. As shown in the column "BIC" in the tables, the BICs for the single beta distributions are larger than that of the two-component beta mixtures by relatively large proportions. And the BICs for the two-component beta mixture distributions are a little larger than that of the three-component beta mixture distributions for both KM and KNM. Obviously, we would prefer the two-component beta mixture distributions, but we would cast a doubt when it comes to the three-component beta mixture distributions. We still prefer the two-component one instead of the three-component one. The reasons are: 1) these differences of BICs between two and three component distributions are really small in proportion (by 0.3% for KM and 0.017% for KNM), 2) the BICs still don't take the sample size effect fully into account since the log function for sample size goes to flat when the sample sizes are large. By simple calculation, we can see that the log-likelihood increased by a multiplicative factor more than 10 when the sample sizes increased from KM case to KNM case, however, at the same time, the log of sample size only roughly increased from 9 to 11. So we choose the two-component beta mixture distributions for both KM and KNM. And also as we have seen, these distributions have potential good forensic interpretations.

It's also helpful to see how the individual components look like in the beta mixture distributions as in Figure 3. The KNM and KM distributions seem to share a common component while keep the other components far apart from each other. The separated components represent the ideal cases where the bullet land engraved areas preserved the information of source well and result in clearly distinct separation. The shared components represent the cases where the KM results in lower scores because of some degree of tank rash, pitting, breakoff or other damages on the bullets and the KNM results in higher score because of the random identification effect. According to the estimated prior probabilities, both distributions put less weight on the common component while put larger weight on the the components characterizing the differences of KM and KNM respectively, which agrees on our expectation that majorities of the distributions are separated while the minorities overlapped. These properties together well explained the observed empirical distribution of similarity scores in our cases. Particularly, the heavier tail of the KM comparisons is explicitly included in the form of the model by one of the components.

Figure 3: Estimated components

Evaluate the error rates

(report theoretical error rates)

Estimations with changing sample size

(quantify the variation, reproduce some results in the previous sections with changing sample size)

Conclusion

References

AFTE Criteria For Identification Committee, 1992. Theory of identification, range striae comparison reports and modified glossary definition afte criteria for identification committee report. AFTE Journal 24, 336–340.

Garton, N., Ommen, D., Niemi, J., Carriquiry, A., 2020. Score-based likelihood ratios to evaluate forensic pattern evidence.

Hare, E., Hofmann, H., Carriquiry, A., 2016. Automatic matching of bullet land impressions. Annals of Applied Statistics 11, 2332–2356.

Hare, E., Hofmann, H., Carriquiry, A., 2017. Algorithmic approaches to match degraded land impressions. Law, Probability and Risk 16, 203–221. doi:10.1093/lpr/mgx018

Krishnan, G., Hofmann, H., 2019. Adapting the chumbley score to match striae on land engraved areas (<scp>LEA</scp> s) of bullets, Journal of Forensic Sciences 64, 728–740. doi:10.1111/1556-4029.13950

National Research Council, 2009. Strengthening forensic science in the united states: A path forward, Strengthening Forensic Science in the United States: A Path Forward. National Academies Press. doi:10.17226/12589

Rice, K., Genschel, U., Hofmann, H., 2020. A robust approach to automatically locating grooves in 3D bullet land scans. Journal of Forensic Sciences 65, 775–783. doi:10.1111/1556-4029.14263

Song, J., Vorburger, T.V., Chu, W., Yen, J., Soons, J.A., Ott, D.B., Zhang, N.F., 2018. Estimating error rates for firearm evidence identifications in forensic science. Forensic Science International 284, 15–32. doi:10.1016/j.forsciint.2017.12.013

Vanderplas, S., Nally, M., Klep, T., Cadevall, C., Hofmann, H., 2020. Comparison of three similarity scores for bullet lea matching. Forensic Science International 308. doi:10.1016/j.forsciint.2020.110167